

CKAD PRACTICE QUESTIONS

Part 1 Questions

1. Kubernetes Architecture

Question:

You are tasked with deploying a multi-container Pod that requires a specific container to always run first before the second one. How would you configure the order of container startup in the Pod specification?

- ☒ A) Use `initContainers` in the Pod specification
 - ☐ B) Use `preStop` hook in the second container
 - ☐ C) Use `readinessProbe` in the first container
 - ☐ D) Use `livenessProbe` in the second container
-

2. Pod Design

Question:

You need to deploy a Pod that runs a web server and logs its output to a file. You also want to ensure that logs are persisted across Pod restarts. How would you accomplish this?

- ☐ A) Mount a `hostPath` volume to the container's log directory
 - ☐ B) Use an `emptyDir` volume to store logs
 - ☒ C) Use a `PersistentVolumeClaim` (PVC) and mount it to the container
 - ☐ D) Use a `ConfigMap` to store the logs
-

3. Configuration and Secrets Management

Question:

You need to store database credentials securely for your application. How would you manage these credentials within Kubernetes?

- ☐ A) Store them in a `ConfigMap`
 - ☒ B) Store them in a `Secret`
 - ☐ C) Store them as environment variables within the Pod
 - ☐ D) Store them as annotations on the Pod
-

4. Networking

Question:

You have two services, `serviceA` and `serviceB`, running in the same namespace. You need to enable communication between them. What is the correct way to allow `serviceA` to communicate with `serviceB`?

- ☒ A) Use a `NetworkPolicy` to allow traffic between the two services
 - ☐ B) Use the DNS name of `serviceB` in the environment variable of `serviceA`
 - ☐ C) Expose both services via LoadBalancer
 - ☐ D) Bind `serviceA` to `serviceB` using a Kubernetes `Pod` definition
-

5. Deployments and Rollouts

Question:

You have a deployment and want to update it with a new image for the application while ensuring zero downtime. Which strategy would you use for this update?

- ☒ A) Use `kubectl apply` with the `--record` flag
 - ☐ B) Use a rolling update strategy with the deployment
 - ☐ C) Use `kubectl rollout restart`
 - ☐ D) Use `kubectl delete` to remove the old pods and create new ones
-

6. Helm Charts

Question:

You are using Helm to deploy an application and need to override a value for a specific configuration. Which command would you use to set a custom value during installation?

- ☒ A) `helm install --set key=value`
 - ☐ B) `helm install --values values.yaml`
 - ☐ C) `helm install --config config.yaml`
 - ☐ D) `helm install --override key=value`
-

7. Pod Scheduling

Question:

You need to ensure that a Pod is scheduled only on nodes that have at least 16 GB of memory. Which of the following methods can you use to achieve this?

- ☐ A) Use a `nodeSelector`
- ☒ B) Use a resource request and limit for memory

- C) Use an affinity rule
 - D) Use a taint and toleration
-

8. Logging and Monitoring

Question:

You need to set up monitoring for your application running in Kubernetes and want to track HTTP requests and response times. Which tool would you use for application-level monitoring?

- A) Prometheus
 - B) Fluentd
 - C) Istio
 - D) Kubernetes Dashboard
-

9. Pod Lifecycle Management

Question:

You have a pod with multiple containers, and you need to execute a command on one of the containers in a running Pod. Which command will allow you to execute the command inside the container?

- A) `kubectl exec <pod-name> --container <container-name> -- <command>`
 - B) `kubectl run <pod-name> --container <container-name> -- <command>`
 - C) `kubectl attach <pod-name> --container <container-name> -- <command>`
 - D) `kubectl exec <container-name> --pod <pod-name> -- <command>`
-

10. Namespaces and Resource Management

Question:

You need to create a new namespace called `dev` and deploy an application that should only run in this namespace. Which command would you use to achieve this?

- A) `kubectl create namespace dev`
 - B) `kubectl set namespace dev`
 - C) `kubectl apply -f app.yaml -n dev`
 - D) `kubectl create deploy app -n dev`
- n does not create a new NS if it does not exist already???

Part 2 Questions

Kubernetes Architecture & Configuration

1. Which of the following statements is true about Kubernetes nodes?
 - A) A node in Kubernetes is a virtual machine only
 - ☒ B) A node can be either a physical or virtual machine
 - C) A node only runs the API server
 - D) Nodes do not run containers in Kubernetes
2. What is the purpose of `kubelet` in a Kubernetes node?
 - A) Manages the Kubernetes cluster's control plane
 - B) Registers nodes with the Kubernetes API server
 - ☒ C) Ensures containers in the node are running as expected
 - D) Provides network services for pods
3. How would you expose a Kubernetes Pod to external traffic?
 - A) Use a Kubernetes Service of type `ClusterIP`
 - B) Use a Kubernetes Service of type `NodePort`
 - C) Use a Kubernetes Service of type `LoadBalancer`
 - ☒ D) Both B and C are correct
4. Which Kubernetes object manages containerized applications' deployment and scaling?
 - ☒ A) ReplicaSet
 - B) Pod
 - C) Deployment
 - D) Service
5. In Kubernetes, how do you define a resource that uses an NFS volume as persistent storage?
 - A) Use `hostPath` volume
 - B) Use `emptyDir` volume
 - ☒ C) Use `nfs` volume
 - D) Use `persistentVolumeClaim` volume

Pod Design & Management

6. Which of the following ensures that a Pod is scheduled to run on a node with at least 4 CPUs?
 - A) `affinity` rules
 - B) `taints` and `tolerations`
 - C) `nodeSelector`
 - ☒ D) Resource requests and limits

7. Which option is used to specify that a container in a Pod should not restart if it fails?
 - A) `restartPolicy: Always`
 - B) `restartPolicy: OnFailure`
 - ☒ C) `restartPolicy: Never`
 - D) `restartPolicy: Delayed`

8. How can you configure a Pod to wait for a specific container to complete before starting the other container?
 - ☒ A) Use `initContainers`
 - B) Use `readinessProbe`
 - C) Use `livenessProbe`
 - D) Use `container.lifecycle` hooks

9. Which of the following is the best way to ensure that logs from a Pod persist across Pod restarts?
 - A) Store logs in a `ConfigMap`
 - B) Mount a `hostPath` volume to the log directory
 - ☒ C) Store logs in a `PersistentVolumeClaim`
 - D) Use `emptyDir` volume

10. You need to define a set of environment variables for your application running in a Pod. Where should you define them?

- A) In the `metadata` section of the Pod
 - ☒ B) In the `spec.containers` section
 - C) In the `spec.volumes` section
 - D) In the `spec.initContainers` section
-

Services, Networking & Communication

11. What is the default type of Kubernetes service?

- ☒ A) `ClusterIP`
- B) `NodePort`
- C) `LoadBalancer`
- D) `ExternalName`

12. How would you allow two Pods in different namespaces to communicate with each other?

- A) Use `NetworkPolicies` to allow cross-namespace traffic
- B) Use `PodSecurityPolicies` to allow cross-namespace traffic
- ☒ C) Use DNS names like `<pod-name>.<namespace>.svc.cluster.local`
- D) Cross-namespace communication is not possible in Kubernetes

13. Which component in Kubernetes manages the DNS for service discovery?

- A) kube-apiserver
- B) kube-proxy
- ☒ C) CoreDNS
- D) etcd

14. What does a `NetworkPolicy` define?

- A) How services communicate with each other
- ☒ B) Which Pods can communicate with each other based on labels and namespaces
- C) Which nodes can access the Kubernetes cluster
- D) How to authenticate users in the cluster

15. How would you ensure that only Pods in a specific namespace can access a service in that namespace?

- ☒ A) Use `NetworkPolicy`
- B) Use `PodSelector`
- C) Use `RBAC`

D) Use `ServiceAccount`

StatefulSet & Persistent Storage

16. What Kubernetes object should you use for managing stateful applications with persistent storage?

- A) Deployment
- ☒ B) StatefulSet
- C) ReplicaSet
- D) DaemonSet

17. What is the difference between `StatefulSet` and `Deployment`?

- ☒ A) `StatefulSet` manages applications that require persistent storage and stable network identities
- B) `StatefulSet` can only manage stateless applications
- C) `StatefulSet` automatically creates `PersistentVolumeClaims` for pods
- D) `StatefulSet` is only for batch processing

18. You have a `StatefulSet` with persistent volumes. What happens when the Pod fails?

- A) A new Pod is created, but the volume is not reused
- ☒ B) A new Pod is created, and the volume is reused
- C) A new Pod is not created automatically
- D) The `PersistentVolume` is deleted

19. Which of the following is the correct way to manage storage in a `StatefulSet`?

- A) Use `emptyDir` for storage
- ☒ B) Use a `PersistentVolumeClaim` in the `volumeClaimTemplates` section
- C) Use a `hostPath` volume
- D) Use `ConfigMap` for storage

20. What is the purpose of a `PersistentVolumeClaim` (PVC)?

- ☒ A) To declare storage requirements for Pods
- B) To bind a Pod to a specific node
- C) To create a new storage resource

D) To manage access control to storage

Deployments & Rollouts

21. How would you roll back a deployment to a previous version in Kubernetes?

- A) Use `kubectl delete deployment <deployment-name>`
- ☒ B) Use `kubectl rollout undo deployment <deployment-name>`
- C) Use `kubectl apply -f <previous-version-file>`
- D) Kubernetes does not support rollbacks

22. Which strategy does Kubernetes use by default when rolling out a new version of a deployment?

- A) Rolling update
- ☒ B) Recreate
- C) Blue-green deployment
- D) Canary deployment

23. Which command allows you to monitor the status of a deployment rollout?

- A) `kubectl get deployment`
- B) `kubectl describe deployment`
- ☒ C) `kubectl rollout status deployment <deployment-name>`
- D) `kubectl get rollout`

24. How can you configure a deployment to run only one replica of a pod?

- ☒ A) Set `replicas: 1` in the deployment manifest
- B) Set `replicas: 0` in the deployment manifest
- C) Use `kubectl scale deployment`
- D) Set `replicas: true`

25. How can you update the image of a running deployment?

- A) `kubectl set image deployment <deployment-name> <container-name>=<new-image>`
- B) Edit the deployment manifest and reapply it
- ☒ C) Both A and B are correct
- D) Delete and recreate the deployment

Helm & Package Management

26. What is the primary purpose of Helm in Kubernetes?

- A) To deploy containers to Kubernetes clusters
- ☒ B) To manage and automate the deployment of applications using charts
- C) To monitor Kubernetes clusters
- D) To control traffic routing in Kubernetes

27. What does the Helm chart `values.yaml` file define?

- A) The Kubernetes cluster's configuration
- ☒ B) The configuration values for the Helm release
- C) The list of available charts
- D) The Helm repository URL

28. Which command would you use to upgrade a Helm release with a new version of a chart?

- ☒ A) `helm upgrade <release-name> <chart-name>`
- B) `helm install <release-name> <chart-name>`
- C) `helm deploy <release-name> <chart-name>`
- D) `helm update <release-name> <chart-name>`

No such subcommand as update

29. How would you install a Helm chart from a local directory?

- ☒ A) `helm install <chart-name> ./chart`
- B) `helm add <chart-name> ./chart`
- C) `helm install ./chart`
- D) `helm install --local <chart-name>`

30. What command would you use to check the values of a Helm release?

- ☒ A) `helm get values <release-name>`
 - B) `helm show values <release-name>`
 - C) `helm values <release-name>`
 - D) `helm status <release-name>`
-

Troubleshooting & Logs

31. Which command would you use to get the logs of a container in a Pod?
- A) `kubectl logs <pod-name>`
 - B) `kubectl logs <container-name> -p`
 - C) `kubectl get logs <pod-name>`
 - D) `kubectl describe logs <pod-name>`
32. If a Pod is stuck in the **ContainerCreating** state, which of the following is a possible reason?
- A) There is an issue with the container image
 - B) The Pod is not scheduled to any node
 - C) The node has insufficient resources for the Pod
 - D) All of the above
33. You want to monitor the resource usage of a Pod. Which command should you use?
- A) `kubectl top pod <pod-name>`
 - B) `kubectl describe pod <pod-name>`
 - C) `kubectl logs <pod-name>`
 - D) `kubectl get pod <pod-name>`
34. What does the **kubectl describe** command provide when run on a Pod?
- A) Pod's logs
 - B) Detailed status and event information about the Pod
 - C) Pod's resource usage
 - D) Pod's containers and their configuration
35. How would you troubleshoot a failing Pod using **kubectl**?
- A) Use `kubectl logs <pod-name>` to check logs
 - B) Use `kubectl describe pod <pod-name>` to check events
 - C) Use `kubectl exec <pod-name> -- <command>` to run debugging commands
 - D) All of the above
-

Security & Access Control

36. Which Kubernetes resource is used to control access to resources within a cluster?

- A) RBAC (Role-Based Access Control)
- B) NetworkPolicy
- ☒ C) ServiceAccount
- D) ConfigMap

37. Which Kubernetes resource defines the permissions for users or service accounts to access resources?

- ☒ A) Role and RoleBinding
- B) PodSecurityPolicy
- C) Deployment
- D) ServiceAccount

38. How can you limit a Pod's access to a specific set of nodes?

- A) Use nodeSelector
- B) Use affinity
- ☒ C) Use taints and tolerations
- D) All of the above

39. Which of the following can be used to run a container with restricted privileges in Kubernetes?

- ☒ A) Use SecurityContext
- B) Use ServiceAccount
- C) Use PodSecurityPolicy
- D) Use NetworkPolicy

40. How can you enforce that containers only run as non-root users in your Kubernetes cluster?

- A) Use PodSecurityPolicy
 - B) Use RBAC
 - ☒ C) Use SecurityContext with runAsUser
 - D) Use NetworkPolicy
-

Other Important Topics

41. What is the function of `kubectl expose` command?
- ☒ A) To create a Kubernetes Service
 - ☐ B) To create a new Deployment
 - ☐ C) To scale a deployment
 - ☐ D) To monitor pod status
42. Which of the following would be the best way to isolate an application that should not be accessed by others in the same cluster?
- ☐ A) Use `NetworkPolicies`
 - ☐ B) Use `PodSecurityPolicies`
 - ☒ C) Use `Namespace`
 - ☐ D) Use `PodSecurityContext`
43. What does the `kubectl scale` command do?
- ☒ A) It resizes the number of Pods in a Deployment, ReplicaSet, or StatefulSet
 - ☐ B) It increases the resource limits of a Pod
 - ☐ C) It increases the number of nodes in the cluster
 - ☐ D) It resizes the volume attached to a Pod
44. You are configuring resource requests and limits for a container in Kubernetes. What happens when a container exceeds its CPU limit?
- ☐ A) The container is restarted automatically
 - ☒ B) The container's CPU usage is throttled
 - ☐ C) The container is evicted
 - ☐ D) The container's memory is expanded
45. What is the purpose of `kubectl apply`?
- ☐ A) To deploy resources to the cluster
 - ☐ B) To delete resources from the cluster
 - ☐ C) To get a list of resources in the cluster
 - ☒ D) To update an existing resource or create it if it does not exist
-

More Questions

46. How do you prevent a Pod from being scheduled on certain nodes?
- A) Use `taints` and `tolerations`
 - B) Use `affinity`
 - ☒ C) Use `nodeSelector`
 - D) All of the above
47. Which command would you use to manually restart a Pod in Kubernetes?
- A) `kubectl restart pod <pod-name>`
 - B) `kubectl delete pod <pod-name>`
 - ☒ C) `kubectl restart deployment <deployment-name>`
 - D) `kubectl rollout restart deployment <deployment-name>`
48. What is the default `kube-proxy` mode in Kubernetes?
- ☒ A) `iptables`
 - B) `ipvs`
 - C) `proxy`
 - D) `loadbalancer`
49. How would you limit the CPU and memory resources for a container in a Pod?
- ☒ A) Use the `limits` and `requests` fields in the container's resource specification
 - B) Use `affinity` rules
 - C) Use a `ServiceAccount`
 - D) Use `NetworkPolicies`
50. What Kubernetes feature allows you to automatically scale the number of Pods in a Deployment?
- A) Horizontal Pod Autoscaler
 - ☒ B) `ReplicaSet`
 - C) `StatefulSet`
 - D) Vertical Pod Autoscaler
-